

WEEK 5

Devesh Mahajan

180905242

B-2

Roll No: 34

1. Write and execute a program in CUDA to add two vectors of length N to meet the following requirements using 3 different kernels

- a) block size as N
- b) N threads within a block
- c) Keep the number of threads per block as 256 (constant) and vary the number of blocks to handle N elements.

```
%%cu

#include <stdio.h>
#include <stdlib.h>
#include <math.h>

__global__ void firstKernel(int *a, int *b, int *c, int n)
{
    int i = blockIdx.x;
    c[i] = a[i] + b[i];
}

__global__ void secondKernel(int *a, int *b, int *c, int n)
{
    int i = threadIdx.x;
    c[i] = a[i] + b[i];
}

__global__ void thirdKernel(int *a, int *b, int *c, int n)
{
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    if (i < n)
        c[i] = a[i] + b[i];
}

int main()
{
    int *d_a, *d_b, *d_c, *d_d, *d_e;
    int n = 10;
    int a[n], b[n];
    for (int i = 0; i < n; i++)
```

```

{
    a[i] = rand() % 50;
    b[i] = rand() % 30;
}
int c[n], d[n], e[n];
printf("First vector is: \n");
for (int i = 0; i < n; i++)
    printf("%d ", a[i]);
printf("\n");
printf("Second vector is: \n");
for (int i = 0; i < n; i++)
    printf("%d ", b[i]);
printf("\n");
cudaMalloc((void **) &d_a, n* sizeof(int));
cudaMalloc((void **) &d_b, n* sizeof(int));
cudaMalloc((void **) &d_c, n* sizeof(int));
cudaMalloc((void **) &d_d, n* sizeof(int));
cudaMalloc((void **) &d_e, n* sizeof(int));
cudaMemcpy(d_a, &a, n* sizeof(int), cudaMemcpyHostToDevice);
cudaMemcpy(d_b, &b, n* sizeof(int), cudaMemcpyHostToDevice);
firstKernel <<<n, 1>>> (d_a, d_b, d_c, n);
secondKernel <<<1, n>>> (d_a, d_b, d_d, n);
thirdKernel <<<ceil(n / 256.0), 256>>> (d_a, d_b, d_e, n);
cudaMemcpy(&c, d_c, n* sizeof(int), cudaMemcpyDeviceToHost);
cudaMemcpy(&d, d_d, n* sizeof(int), cudaMemcpyDeviceToHost);
cudaMemcpy(&e, d_e, n* sizeof(int), cudaMemcpyDeviceToHost);
printf("\n");
printf("The sum using N blocks is: \n");
for (int i = 0; i < n; i++)
    printf("%d ", c[i]);
printf("\n\n");
printf("The sum using N threads is: \n");
for (int i = 0; i < n; i++)
    printf("%d ", d[i]);
printf("\n\n");
printf("The sum using 256 threads is: \n");
for (int i = 0; i < n; i++)
    printf("%d ", e[i]);
printf("\n\n");
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);
cudaFree(d_d);
cudaFree(d_e);
return 0;
}

```

```
First vector is:  
33 27 43 36 49 12 40 13 40 22  
Second vector is:  
16 25 25 12 1 7 19 16 6 16
```

```
The sum using N blocks is:  
49 52 68 48 50 19 59 29 46 38
```

```
The sum using N threads is:  
49 52 68 48 50 19 59 29 46 38
```

```
The sum using 256 threads is:  
49 52 68 48 50 19 59 29 46 38
```

2. Write and execute a CUDA program to read an array of N integer values. Sort the array in parallel using parallel selection sort and store the result in another array.

```
%%cu  
  
#include <stdio.h>  
#include <stdlib.h>  
  
__global__ void parallelKernel(int *a, int n)  
{  
    int i = threadIdx.x;  
    int data = a[i];  
    int pos = 0;  
    for (int j = 0; j < n; j++)  
    {  
        if (a[j] < data || (a[j] == data && j < i))  
            pos++;  
    }  
  
    a[pos] = data;  
}  
  
void sortArr(int* h_arr, int n){  
    int* d_arr = NULL;  
    int size = n*sizeof(int);  
    cudaError_t err = cudaSuccess;  
    err = cudaMalloc((void **) &d_arr, size);  
    if(err != cudaSuccess){  
        printf("%s\n",cudaGetErrorString(err));  
        exit(EXIT_FAILURE);  
    }
```

```

err = cudaMemcpy(d_arr,h_arr,size,cudaMemcpyHostToDevice);
if(err != cudaSuccess){
    printf("%s\n",cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}
parallelKernel<<<1,n>>>(d_arr,n);
cudaMemcpy(h_arr,d_arr,size,cudaMemcpyDeviceToHost);
cudaFree(d_arr);
}

int main() {
    int n = 10;
    int* arr = (int*)malloc(n*sizeof(int));
    for(int i=0;i<n;i++){
        arr[i] = rand()%50;
    }
    printf("Entered Array: \n");
    for(int i=0;i<10;i++){
        printf("%d ",arr[i]);
    }
    printf("\nSorted Array: \n");
    sortArr(arr,n);
    for(int i=0;i<10;i++){
        printf("%d ",arr[i] );
    }
    return 0;
}

```

→ Entered Array:
43 46 57 55 53 55 46 12 9 1
Sorted Array:
1 9 12 43 46 46 53 55 55 57

3. Write and execute a CUDA program to read an integer array of size N. Sort this array using odd-even transposition sorting. Use 2 kernels.

```

%%cu

#include <stdio.h>
#include <stdlib.h>
#include <cuda.h>
#include <cuda_runtime.h>

__global__ void oddKernel(int* arr, int n){


```

```

int i = blockIdx.x*blockDim.x+threadIdx.x;
if(i<n-1 && i%2!=0){
    if(arr[i] > arr[i+1]){
        int temp = arr[i];
        arr[i] = arr[i+1];
        arr[i+1] = temp;
    }
}
}

global__ void evenKernel(int* arr, int n){
int i = blockIdx.x*blockDim.x+threadIdx.x;
if(i<n-1 && i%2==0){
    if(arr[i] > arr[i+1]){
        int temp = arr[i];
        arr[i] = arr[i+1];
        arr[i+1] = temp;
    }
}
}

void oddEvenSort(int* h_arr, int n){
int* d_arr = NULL;
int size = n*sizeof(int);
cudaError_t err = cudaSuccess;
err = cudaMalloc((void **) &d_arr, size);
if(err != cudaSuccess){
    printf("%s\n", cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}
err = cudaMemcpy(d_arr, h_arr, size, cudaMemcpyHostToDevice);
if(err != cudaSuccess){
    printf("%s\n", cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}
int i = 0;
while(i<=n/2) {
    oddKernel<<<1,n>>>(d_arr,n);
    evenKernel<<<1,n>>>(d_arr,n);
    i++;
}
cudaMemcpy(h_arr, d_arr, size, cudaMemcpyDeviceToHost);
cudaFree(d_arr);
}

int main(){
    int n = 10;
    int* arr = (int*)malloc(n*sizeof(int));
    for(int i=0;i<n;i++){
        arr[i] = rand()%50;
    }
}

```

```
printf("Entered Array: \n");
for(int i=0;i<10;i++){
    printf("%d ",arr[i]);
}
printf("\nSorted Array: \n");
oddEvenSort(arr,n);
for(int i=0;i<10;i++){
    printf("%d ",arr[i] );
}
return 0;
}
```

↳ Entered Array:
33 36 27 15 43 35 36 42 49 21
Sorted Array:
15 21 27 33 35 36 36 42 43 49