

WEEK9 : Programs on Matrix using CUDA (Part1)

Name: Sagnik Chatterjee

Reg: 180905478

Roll No: 61

SEC: VIB

Q1 Write a program in CUDA to add two Matrices for the following specifications:

- a. Each row of resultant matrix to be computed by one thread.**
- b. Each column of resultant matrix to be computed by one thread.**
- c. Each element of resultant matrix to be computed by one thread.**

```
% % cu
#include <stdio.h>
#include <stdlib.h>
#include <math.h>

//each row computed by one thread : - part1
__global__ void row_one_thread(int *a, int *b, int *c, int m,
int n)
{
    int idx = threadIdx.x;
    printf("idx = %d\n", idx);
    for (int i = 0; i < m; ++i)
    {
        c[i + m * idx] = a[i + m * idx] + b[i + m * idx];
    }
}

//each col computed by one thread :- part2
__global__ void col_one_thread(int *a, int *b, int *c, int m,
int n)
{
    int idx = threadIdx.x;
    printf("idx = %d\n", idx);
```

```

for (int i = 0; i < n; ++i, idx += m)
{
    c[idx] = a[idx] + b[idx];
}

//each element computed by one thread :- part3
__global__ void element_one_thread(int *a, int *b, int *c)
{
    int idx = blockDim.x * blockIdx.x + threadIdx.x;
    printf("idx = %d\n", idx);
    c[idx] = a[idx] + b[idx];
}

int main(void)
{
    int i, j;
    int *d_a, *d_b, *d_c, *d_d, *d_e;
    int m = 3;
    int n = 4;
    int a[m * n], b[m * n], c[m * n], d[m * n], e[m * n];

    for (i = 0; i < m * n; ++i)
    {
        a[i] = i + 1;
        b[i] = m * n - i - 1;
    }

    int size = sizeof(int) * m * n;

    cudaMalloc((void **) &d_a, size);
    cudaMalloc((void **) &d_b, size);
    cudaMalloc((void **) &d_c, size);
    cudaMalloc((void **) &d_d, size);
}

```

```

cudaMalloc( (void **) &d_e, size);

printf("Matrix a:\n");
for (i = 0; i < n; ++i)
{
    for (j = 0; j < m; ++j)
    {
        printf("%d\t", a[i * m + j]);
    }
    printf("\n");
}
printf("\n");

printf("Matrix B:\n");
for (i = 0; i < n; ++i)
{
    for (j = 0; j < m; ++j)
    {
        printf("%d\t", b[i * m + j]);
    }
    printf("\n");
}
printf("\n");

// Copy inputs to device
cudaMemcpy(d_a, &a, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, &b, size, cudaMemcpyHostToDevice);

row_one_thread<<<1, n>>>(d_a, d_b, d_c, m, n);
col_one_thread<<<1, n>>>(d_a, d_b, d_d, m, n);
element_one_thread<<<m, n>>>(d_a, d_b, d_e);

cudaMemcpy(&c, d_c, size, cudaMemcpyDeviceToHost);
cudaMemcpy(&d, d_d, size, cudaMemcpyDeviceToHost);

```

```
cudaMemcpy(&e, d_e, size, cudaMemcpyDeviceToHost);

printf("Part 1> A+B:\n");
for (i = 0; i < n; ++i)
{
    for (j = 0; j < m; ++j)
    {
        printf("%d\t", c[i * m + j]);
    }
    printf("\n");
}
printf("\n");

printf("Part 2> A+B:\n");
for (i = 0; i < n; ++i)
{
    for (j = 0; j < m; ++j)
    {
        printf("%d\t", d[i * m + j]);
    }
    printf("\n");
}
printf("\n");

printf("Part 3> A+B:\n");
for (i = 0; i < n; ++i)
{
    for (j = 0; j < m; ++j)
    {
        printf("%d\t", e[i * m + j]);
    }
    printf("\n");
}
printf("\n");
```

```
    cudaFree(d_a);  
    cudaFree(d_b);  
    cudaFree(d_c);  
    cudaFree(d_d);  
    cudaFree(d_e);  
  
    return 0;  
}
```

Output:

Q2 Write a program in CUDA to multiply two Matrices for the following specifications:

- a. Each row of resultant matrix to be computed by one thread.
 - b. Each column of resultant matrix to be computed by one thread.
 - c. Each element of resultant matrix to be computed by one thread.

```
%%cu  
#include <stdio.h>  
#include <stdlib.h>
```

```

#include <math.h>

//each row computed by one thread : - part1
__global__ void row_one_thread (int *a, int *b, int *c,int m
,int n) {
    int idx = threadIdx.x;
    printf("idx = %d\n", idx);
    for (int i = 0; i < m; ++i) {
        c[i + m * idx] = a[i + m * idx] * b[i + m * idx];
    }
}

//each col computed by one thread :- part2
__global__ void col_one_thread (int *a, int *b, int *c,int m
,int n ) {
    int idx = threadIdx.x;
    printf("idx = %d\n", idx);
    for (int i = 0; i < n; ++i, idx += m) {
        c[idx] = a[idx] * b[idx];
    }
}

//each element computed by one thread :- part3
__global__ void element_one_thread (int *a, int *b, int *c ) {
    int idx = blockDim.x * blockIdx.x + threadIdx.x;
    printf("idx = %d\n", idx);
    c[idx] = a[idx] * b[idx];
}

int main(void) {
    int i,j;
    int *d_a,*d_b,*d_c,*d_d,*d_e;
    int m =3;

```

```
int n =4 ;
int a[m*n],b[m*n],c[m*n],d[m*n],e[m*n] ;

for (i = 0; i < m * n; ++i) {
    a[i] = i + 1;
    b[i] = m * n - i - 1;
}

int size = sizeof(int) * m * n;

cudaMalloc((void**) &d_a, size);
cudaMalloc((void**) &d_b, size);
cudaMalloc((void**) &d_c, size);
cudaMalloc((void**) &d_d, size);
cudaMalloc((void**) &d_e, size);

printf("Matrix a:\n");
for (i = 0; i < n; ++i) {
    for (j = 0; j < m; ++j) {
        printf("%d\t", a[i * m + j]);
    }
    printf("\n");
}
printf("\n");

printf("Matrix B:\n");
for (i = 0; i < n; ++i) {
    for (j = 0; j < m; ++j) {
        printf("%d\t", b[i * m + j]);
    }
    printf("\n");
}
printf("\n");
```

```

// Copy inputs to device
cudaMemcpy(d_a, &a, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_b, &b, size, cudaMemcpyHostToDevice);

row_one_thread<<<1,n>>>(d_a,d_b,d_c,m,n);
col_one_thread<<<1, n>>>(d_a, d_b, d_d,m,n);
element_one_thread<<<m,n>>>(d_a,d_b,d_e);

cudaMemcpy(&c, d_c, size, cudaMemcpyDeviceToHost);
cudaMemcpy(&d, d_d, size, cudaMemcpyDeviceToHost);
cudaMemcpy(&e, d_e, size, cudaMemcpyDeviceToHost);

printf("Part 1> A*B:\n");
for (i = 0; i < n; ++i) {
    for (j = 0; j < m; ++j) {
        printf("%d\t", c[i * m + j]);
    }
    printf("\n");
}
printf("\n");

printf("Part 2> A*B:\n");
for (i = 0; i < n; ++i) {
    for (j = 0; j < m; ++j) {
        printf("%d\t", d[i * m + j]);
    }
    printf("\n");
}
printf("\n");

printf("Part 3> A*B:\n");
for (i = 0; i < n; ++i) {
    for (j = 0; j < m; ++j) {

```

```
    printf("%d\t", e[i * m + j]);
}
printf("\n");
printf("\n");

cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);
cudaFree(d_d);
cudaFree(d_e);

return 0;
}
```

Output: